

Master UI with 2 workers : 192.168.188.21 local machine, 192.168.88.60 Raspberry Pi

The screenshot shows the Spark Master UI interface. At the top, there are tabs for "E-Mail – Raichle, Lorena – Outlook", "localhost", "Spark Master at spark://192.168.188.21:7077", "Spark shell - Spark Jobs", and "Browsing HDFS". Below the tabs, the title is "Spark Master at spark://192.168.188.21:7077". The main content area displays the following statistics:

- URL: spark://192.168.188.21:7077
- Alive Workers: 2
- Cores in use: 12 Total, 0 Used
- Memory in use: 9.7 GiB Total, 0.0 B Used
- Resources in use:
- Applications: 0 Running, 7 Completed
- Drivers: 0 Running, 0 Completed
- Status: ALIVE

Workers (2)

Worker Id	Address	State	Cores	Memory	Resources
worker-20231107114650-192.168.188.21:49279	192.168.188.21:49279	ALIVE	8 (0 Used)	7.0 GiB (0.0 B Used)	
worker-20231107115116-192.168.188.60:45479	192.168.188.60:45479	ALIVE	4 (0 Used)	2.7 GiB (0.0 B Used)	

Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration

Completed Applications (7)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
app-20231107124655-0006	WordCountWithSpark	12	1024.0 MiB		2023/11/07 12:46:55	Lorena	FINISHED	7 s
app-20231107123718-0005	WordCountWithSpark	12	1024.0 MiB		2023/11/07 12:37:18	Lorena	FINISHED	2 s
app-20231107120907-0004	WordCountWithSpark	12	1024.0 MiB		2023/11/07 12:09:07	Lorena	FINISHED	2 s
app-20231107120709-0003	WordCountWithSpark	12	1024.0 MiB		2023/11/07 12:07:09	Lorena	FINISHED	6 s
app-20231107120351-0002	WordCountWithSpark	12	1024.0 MiB		2023/11/07 12:03:51	Lorena	FINISHED	6 s
app-20231107120020-0001	WordCountWithSpark	12	1024.0 MiB		2023/11/07 12:00:20	Lorena	FINISHED	6 s
app-20231107115228-0000	WordCountWithSpark	12	1024.0 MiB		2023/11/07 11:52:28	Lorena	FINISHED	7 s

App-20231107123718-006

The screenshot shows the Application UI for "WordCountWithSpark". At the top, there are tabs for "E-Mail – Raichle, Lorena – Outlook", "localhost", "Application: WordCountWithSpark", "Spark shell - Spark Jobs", and "Browsing HDFS". Below the tabs, the title is "Application: WordCountWithSpark". The main content area displays the following application details:

- ID: app-20231107124655-0006
- Name: WordCountWithSpark
- User: Lorena
- Cores: Unlimited (12 granted)
- Executor Limit: Unlimited (2 granted)
- Executor Memory - Default Resource Profile: 1024.0 MiB
- Executor Resources - Default Resource Profile:
- Submit Date: 2023/11/07 12:46:55
- Status: FINISHED

Executor Summary (2)

ExecutorID	Worker	Cores	Memory	Resource Profile Id	Resources	State	Logs

Removed Executors (2)

ExecutorID	Worker	Cores	Memory	Resource Profile Id	Resources	State	Logs
1	worker-20231107115116-192.168.188.60-45479	4	1024	0		KILLED	stdout stderr
0	worker-20231107114650-192.168.188.21-49279	8	1024	0		KILLED	stdout stderr

Executor ID 0

The screenshot shows the Executor UI for Executor ID 0. At the top, there are tabs for "E-Mail – Raichle, Lorena – Outlook", "localhost", "Spark Worker at 192.168.188.21:49279", "Spark shell - Spark Jobs", and "Browsing HDFS". Below the tabs, the title is "Spark Worker at 192.168.188.21:49279". The main content area displays the following executor details:

- ID: worker-20231107114650-192.168.188.21-49279
- Master URL: spark://192.168.188.21:7077
- Cores: 8 (0 Used)
- Memory: 7.0 GiB (0.0 B Used)
- Resources:

[Back to Master](#)

Running Executors (0)

ExecutorID	State	Cores	Memory	Resources	Job Details	Logs

Finished Executors (7)

ExecutorID	State	Cores	Memory	Resources	Job Details	Logs
0	KILLED	8	1024.0 MiB		ID: app-20231107124655-0006 Name: WordCountWithSpark User: Lorena	stdout stderr

Executor ID 1

Spark Worker at 192.168.188.60:45479

ID: worker-2023110715116-192.168.188.60-45479
 Master URL: spark://192.168.188.21:7077
 Cores: 4 (0 Used)
 Memory: 2.7 GB (0.0 B Used)
 Resources:

[Back to Master](#)

Running Executors (0)

ExecutorID	State	Cores	Memory	Resources	Job Details	Logs
------------	-------	-------	--------	-----------	-------------	------

Finished Executors (7)

ExecutorID	State	Cores	Memory	ID: app-20231107124655-0006 Name: WordCountWithSpark User: Lorena	stdout stderr
1	KILLED	4	1024.0 MiB		

Localhost:8080 mac (worker)

Apache Spark 3.5.0 Jobs Stages Storage Environment Executors

Spark Jobs (?)

User: Lorena
 Total Uptime: 1,0 h
 Scheduling Mode: FIFO

Event Timeline Enable zooming

The timeline chart shows two main sections: Executors and Jobs. The Executors section has a legend for 'Added' (blue) and 'Removed' (red). The Jobs section has a legend for 'Succeeded' (blue), 'Failed' (red), and 'Running' (green). A vertical blue line marks the event 'Executor driver added' at approximately 22 minutes on November 7, 2023. The x-axis is labeled with 15, 20, 25, and 30.

7 November 11:48

HDFS: localhost:8080

localhost

E-Mail – Raichie, Lorena – Outlook | Spark Worker at 192.168.188.60:45479 | Spark shell - Spark Jobs | Browsing HDFS

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Browse Directory

/sparkApp

Show 25 entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
<input type="checkbox"/>	drwxr-xr-x	Lorena	supergroup	0 B	Nov 07 12:46	0	0 B	output_task2_final
<input type="checkbox"/>	drwxr-xr-x	Lorena	supergroup	0 B	Nov 07 12:07	0	0 B	output_12
<input type="checkbox"/>	drwxr-xr-x	Lorena	supergroup	0 B	Nov 07 12:03	0	0 B	output_11
<input type="checkbox"/>	drwxr-xr-x	Lorena	supergroup	0 B	Nov 07 12:00	0	0 B	output_10
<input type="checkbox"/>	drwxr-xr-x	Lorena	supergroup	0 B	Nov 07 11:52	0	0 B	output_9
<input type="checkbox"/>	drwxr-xr-x	Lorena	supergroup	0 B	Nov 06 21:40	0	0 B	output_7
<input type="checkbox"/>	drwxr-xr-x	Lorena	supergroup	0 B	Nov 06 21:32	0	0 B	output_6
<input type="checkbox"/>	drwxr-xr-x	Lorena	supergroup	0 B	Nov 06 21:15	0	0 B	output_5

Browse Directory

/sparkApp/output_task2_final Go!

Show 25 entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	<input type="checkbox"/>
<input type="checkbox"/>	-rw-r--r--	Lorena	supergroup	0 B	Nov 07 12:47	3	128 MB	_SUCCESS	<input type="checkbox"/>
<input type="checkbox"/>	-rw-r--r--	Lorena	supergroup	1.92 KB	Nov 07 12:47	3	128 MB	part-00000	<input type="checkbox"/>
<input type="checkbox"/>	-rw-r--r--	Lorena	supergroup	1.99 KB	Nov 07 12:47	3	128 MB	part-00001	<input type="checkbox"/>

Showing 1 to 3 of 3 entries Previous 1 Next

Hadoop, 2021.

File information - part-00001

[Download](#) [Head the file \(first 32K\)](#) [Tail the file \(last 32K\)](#)

Block information -- Block 0

Block ID: 1073741862
 Block Pool ID: BP-83492355-192.168.188.21-1699270201992
 Generation Stamp: 1038
 Size: 2041
 Availability:
 • mbp-lorena.fritz.box

File contents

```
(end,33)
(houses,1)
(chuckle,1)
(don't,1)
(patience,1)
(safe,1)
(doubts,1)
```

File information - part-00000

[Download](#) [Head the file \(first 32K\)](#) [Tail the file \(last 32K\)](#)

Block information -- Block 0

Block ID: 1073741861
 Block Pool ID: BP-83492355-192.168.188.21-1699270201992
 Generation Stamp: 1037
 Size: 1963
 Availability:
 • mbp-lorena.fritz.box

File contents

```
(pain,1)
(its,1)
(staving,1)
(left,1)
(have,1)
(excited,1)
(been,1)
(who,2)
```